

# PAR Laboratory Assignment

## Lab 4: Divide and Conquer parallelism with OpenMP: Sorting

Group 13-03

Carlos Bárcenas Holguera

Miquel Gómez i Esteve

Facultat d'Informàtica de Barcelona - UPC

March 2018

Spring 2017-2018

## Index

### Part I:

**"Divide and conquer" .....3**

### Part II:

**Shared-memory parallelization with OpenMP tasks....6**

### Part III:

**Using OpenMP task dependencies.....11**

## Part I: "Divide and conquer"

We will try to analyse the program with tareador. The first thing is to do is to add the the tareador API's calls to the code as shown in figure 1.

```
void merge(long n, T left[n], T right[n], T result[n*2], long start, long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        tareador_start_task("merge");
        merge(n, left, right, result, start, length/2);
        tareador_end_task("merge");
        tareador_start_task("merge");
        merge(n, left, right, result, start + length/2, length/2);
        tareador_end_task("merge");
    }
}

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition

        tareador_start_task("multisort");
        multisort(n/4L, &data[0], &tmp[0]);
        tareador_end_task("multisort");

        tareador_start_task("multisort");
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);
        tareador_end_task("multisort");

        tareador_start_task("multisort");
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);
        tareador_end_task("multisort");

        tareador_start_task("multisort");
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);
        tareador_end_task("multisort");

        tareador_start_task("merge");
        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
        tareador_end_task("merge");

        tareador_start_task("merge");
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);
        tareador_end_task("merge");

        tareador_start_task("merge");
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
        tareador_end_task("merge");
    } else {
        // Base case
        basicsort(n, data);
    }
}
```

Figure 1: multisort-tareador.c with the tareador's calls added.

Once we have modified the code we can execute the program and we can obtain the dependence graph from figure 2.




Figure 2:  
Task dependence graph

We can see that in every case the vector is divided in four segments, until the segments obtained are smaller than the preset minimum size, which is  $\text{MIN\_SORT\_SIZE} * 4L$ . Once the segments are smaller, they are ordered on the base case, with the basicsort function.

After that we enter the merge part of the code where we have a minimum size for us to execute the merge function, which is  $\text{MIN\_MERGE\_SIZE} * 2L$ . If the segments are smaller we merge them with the basicmerge function.

All the dependencies shown in the graph are caused by the fact that we need the segments to be ordered before we merge them.

We will now use the Tareador's predictions to see how our program would behave. Figure 3 are the results obtained with different number of threads.

| # threads | Execution Time | SpeedUp |
|-----------|----------------|---------|
| 1         | 20.334.421     | -       |
| 2         | 10.173.740     | 1,998   |
| 4         | 5.087.595      | 3,997   |
| 8         | 2.547.419      | 7,983   |
| 16        | 1.289.941      | 15,775  |
| 32        | 1.289.920      | 15,779  |
| 64        | 1.289.909      | 15,779  |

Figure 3: Times obtained with Tareador predictions

From the results obtained, we can prove that once we have 16 threads, we have reached the maximum parallelism of the program. This means that we cannot have a higher speedup and any other threads we want to add will not be used in the execution.

## Part II: Shared-memory parallelization with OpenMP tasks

We will now analyse the program parallelization and performance with tasks. First of all we must modify the code to implement the two versions: leaf and tree. The figures 4 and 5 show the modified codes for the leaf and tree versions respectively.

```
void merge(long n, T left[n], T right[n], T result[n*2], long start, long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        #pragma omp task
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        merge(n, left, right, result, start, length/2);
        merge(n, left, right, result, start + length/2, length/2);
    }
}

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        multisort(n/4L, &data[0], &tmp[0]);
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);
        #pragma omp taskwait

        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);
        #pragma omp taskwait

        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
    } else {
        // Base case
        #pragma omp task
        basicsort(n, data);
    }
}
```

Figure 4: Leaf Strategy Code

```

void merge(long n, T left[n], T right[n], T result[n*2], long start, long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        #pragma omp task
        merge(n, left, right, result, start, length/2);
        #pragma omp task
        merge(n, left, right, result, start + length/2, length/2);
        #pragma omp taskwait
    }
}

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        #pragma omp task
        multisort(n/4L, &data[0], &tmp[0]);
        #pragma omp task
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);
        #pragma omp task
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);
        #pragma omp task
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);
        #pragma omp taskwait

        #pragma omp task
        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
        #pragma omp task
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);
        #pragma omp taskwait

        #pragma omp task
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
        #pragma omp taskwait
    } else {
        // Base case
        basicsort(n, data);
    }
}

```

Figure 5: Tree Strategy Code

In the leaf strategy, all we add is a openmp task construct every time we use the functions basicsort and basicmerge. However, in the tree strategy we need to add the task construct every time we use the function multisort or merge, and for the program to work properly, some taskwait constructs are needed as well.

We will now see how the leaf version behaves when checking its strong scalability.




Figure 6:  
Plots for Leaf Strategy





Figure 7:  
Leaf Strategy Trace  
with 8 threads



|              | Running             | Not created      | Synchronization | Scheduling and Fork/Join | I/O           |
|--------------|---------------------|------------------|-----------------|--------------------------|---------------|
| THREAD 1.1.1 | 1,900,632,297 ns    | -                | -               | 1,051,065,678 ns         | 657,969 ns    |
| THREAD 1.1.2 | 734,808,130 ns      | 1,015,305,440 ns | -               | 918 ns                   | 768,394 ns    |
| THREAD 1.1.3 | 652,533,255 ns      | 1,015,306,012 ns | -               | 985 ns                   | 727,494 ns    |
| THREAD 1.1.4 | 730,750,578 ns      | 1,015,304,287 ns | -               | 923 ns                   | 726,560 ns    |
| THREAD 1.1.5 | 853,195,884 ns      | 1,015,308,445 ns | -               | 870 ns                   | 757,465 ns    |
| THREAD 1.1.6 | 719,048,426 ns      | 1,015,210,682 ns | -               | 5,195 ns                 | 72,268 ns     |
| THREAD 1.1.7 | 774,297,852 ns      | 1,015,322,475 ns | -               | 793 ns                   | 53,651 ns     |
| THREAD 1.1.8 | 1,660,986,151 ns    | 1,015,210,674 ns | 200,724,093 ns  | 9,036,168 ns             | 875,662 ns    |
| Total        | 8,026,252,573 ns    | 7,106,968,015 ns | 200,724,093 ns  | 1,060,111,530 ns         | 4,639,463 ns  |
| Average      | 1,003,281,571.62 ns | 1,015,281,145 ns | 200,724,093 ns  | 132,513,941.25 ns        | 579,932.88 ns |
| Maximum      | 1,900,632,297 ns    | 1,015,322,475 ns | 200,724,093 ns  | 1,051,065,678 ns         | 875,662 ns    |
| Minimum      | 652,533,255 ns      | 1,015,210,674 ns | 200,724,093 ns  | 793 ns                   | 53,651 ns     |

Figure 8:  
Leaf Strategy Statistics  
with 8 threads

As we can see in figures 6, 7 and 8, the leaf strategy does not behave very well so its scalability is not strong at all. The reason must be that we must wait until the recursive part finishes so that we can create the different tasks and distribute them between the rest of the threads. At the end we also have a lot of synchronizations that will affect negatively on the execution time.

It's time to now check the behavior of the tree version.




Figure 9:  
Plots for Tree Strategy





Figure 10:  
Tree Strategy Trace  
with 8 threads



|              | Running          | Not created         | Synchronization | Scheduling and Fork/Join | I/O           |
|--------------|------------------|---------------------|-----------------|--------------------------|---------------|
| THREAD 1.1.1 | 1,890,107,176 ns | -                   | 89,979,365 ns   | 10,789,387 ns            | 355,488 ns    |
| THREAD 1.1.2 | 857,461,434 ns   | 1,026,059,469 ns    | 25,182,480 ns   | 1,040,825 ns             | 279,916 ns    |
| THREAD 1.1.3 | 880,420,941 ns   | 1,026,058,727 ns    | 15,439,666 ns   | 922,426 ns               | 310,645 ns    |
| THREAD 1.1.4 | 862,892,409 ns   | 1,026,059,409 ns    | 19,277,139 ns   | 891,938 ns               | 277,739 ns    |
| THREAD 1.1.5 | 794,327,567 ns   | 1,026,054,429 ns    | 103,044,033 ns  | 751,521 ns               | 258,748 ns    |
| THREAD 1.1.6 | 591,043,723 ns   | 1,025,964,133 ns    | 291,624,947 ns  | 712,513 ns               | 187,508 ns    |
| THREAD 1.1.7 | 858,660,048 ns   | 1,026,069,950 ns    | 32,680,945 ns   | 782,323 ns               | 283,092 ns    |
| THREAD 1.1.8 | 782,780,246 ns   | 1,025,963,355 ns    | 115,967,577 ns  | 614,486 ns               | 279,297 ns    |
| Total        | 7,517,693,544 ns | 7,182,229,472 ns    | 693,176,152 ns  | 16,505,419 ns            | 2,232,433 ns  |
| Average      | 939,711,693 ns   | 1,026,032,781.71 ns | 86,647,019 ns   | 2,063,177.38 ns          | 279,054.12 ns |
| Maximum      | 1,890,107,176 ns | 1,026,069,950 ns    | 291,624,947 ns  | 10,789,387 ns            | 355,488 ns    |
| Minimum      | 591,043,723 ns   | 1,025,963,355 ns    | 15,439,666 ns   | 614,486 ns               | 187,508 ns    |

Figure 11:  
Tree Strategy Statistics  
with 8 threads

The behaviour and scalability of the tree version is a much better as we can see in figures 9, 10 and 11. This is due to the fact that we create tasks during the recursive part of the program, so that its execution is much faster, while in the leaf version tasks were only created in the base cases. We do not have that many synchronizations because of the taskwait construct we added in several places, and we can assure that the data is already ordered before getting to the merge function for example.

Comparing the speed-up plot for the multisort execution with the ideal speed-up values obtained with tareador, we can assure that this is a very good parallelization since they are very similar in most of the values for #threads, and also because the tareador values are ideal and do not check any of the overheads.

However there is an exception on the last execution with 12 threads.



## Optional:

When submitting for execution with 8 threads we see that the initialization part without any kind of parallelization of the program takes 0.797846 seconds and the trace resulting would be the same shown in figure 10 where is clearly represented the initialization part, which uses only one thread.

Next, we parallelize the functions used to initialize the vectors used in the program as seen the the next figure (12).

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    initialize(N, data);
    #pragma omp task
    clear(N, tmp);
    #pragma omp taskwait
}
```

Figure 12: Modified code to parallelize program initialization

When submitting for execution with 8 threads we see a very little improvement since the initialization time is 0.753048 seconds. The total speed-up is  $0.753048/0.797846 = 1,06$ . The reason behind this poor speed-up would be the disbalance between the two tasks, as we can see in figure 14, where the thread 1 executes the function initialize while the thread 4 executes the clear function, which takes very little time compared to the initialize function.




Figure 13:  
Tree Strategy Trace  
and parallelized  
initialization with 8  
threads.

After analyzing the code we can see that the function initialize has a clear dependency (to calculate  $\text{data}[i]$  we need the value of  $\text{data}[i-1]$ ) which forces the execution to be sequential.

Because of that we think it wouldn't make sense to try to parallelize the two functions more efficiently (with a parallel for) since the improvement margin is very small.

## Part III: Using OpenMP task dependencies

First of all we need to change the code adding all the in, out and inout dependencies of the different tasks of the code. The changed code is shown in figure 15

```
void merge(long n, T left[n], T right[n], T result[n*2], long start, long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        #pragma omp task
        merge(n, left, right, result, start, length/2);
        #pragma omp task
        merge(n, left, right, result, start + length/2, length/2);
        #pragma omp taskwait
    }
}

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        #pragma omp task depend(inout: data[0])
        multisort(n/4L, &data[0], &tmp[0]);

        #pragma omp task depend(inout: data[n/4L])
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);

        #pragma omp task depend(inout: data[n/2L])
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);

        #pragma omp task depend(inout: data[3L*n/4L])
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);

        #pragma omp task depend(in: data[0], data[n/4L]) depend(out: tmp[0])
        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
        #pragma omp task depend(in: data[n/2L], data[3L*n/4L]) depend(out: tmp[n/2L])
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);

        #pragma omp task depend(in: tmp[0], tmp[n/2L]) depend(out: data[0])
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);

        #pragma omp taskwait
    } else {
        // Base case
        basicsort(n, data);
    }
}
```

Figure 14: Modified code to control dependencies between tasks

Then we submit the program for execution with 8 threads and the execution time obtained for the multisort is 0.996871 seconds.

It's now time to check the strong scalability of our program, so we submit for execution the code with the submit-strong script. The results obtained as plots are shown in figure 16, and we can see that it is very similar to the last version analyzed (tree version with taskwait to ensure dependencies are satisfied)




Figure 15: Speed Up plot for the task dependencies version

The speed-up obtained in the different cases is pretty close to the ideal one given by the taredador. We can see that it has a very similar scalability than the one shown in the tree task version, mostly because even though the dependencies declarations are different, the dependencies graph would be nearly the same except for some minor differences (We now know that all the dependencies generated are necessary).




Figure 16: Trace for task dependencies version

|                | Running  | Not created | Synchronization | Scheduling and Fork/Join | I/O    | Others |
|----------------|----------|-------------|-----------------|--------------------------|--------|--------|
| THREAD 1.1.1   | 94.97 %  | -           | 4.75 %          | 0.27 %                   | 0.01 % | 0.00 % |
| THREAD 1.1.2   | 44.36 %  | 51.33 %     | 4.25 %          | 0.06 %                   | 0.01 % | -      |
| THREAD 1.1.3   | 47.70 %  | 51.32 %     | 0.92 %          | 0.05 %                   | 0.01 % | -      |
| THREAD 1.1.4   | 48.63 %  | 51.25 %     | 0.05 %          | 0.06 %                   | 0.01 % | -      |
| THREAD 1.1.5   | 44.91 %  | 51.27 %     | 3.75 %          | 0.07 %                   | 0.01 % | -      |
| THREAD 1.1.6   | 0.00 %   | 51.21 %     | 48.78 %         | 0.01 %                   | 0.00 % | -      |
| THREAD 1.1.7   | 37.12 %  | 51.21 %     | 11.60 %         | 0.06 %                   | 0.01 % | -      |
| THREAD 1.1.8   | 48.21 %  | 51.32 %     | 0.40 %          | 0.06 %                   | 0.01 % | -      |
|                |          |             |                 |                          |        |        |
| <b>Total</b>   | 365.89 % | 358.91 %    | 74.49 %         | 0.63 %                   | 0.08 % | 0.00 % |
| <b>Average</b> | 45.74 %  | 51.27 %     | 9.31 %          | 0.08 %                   | 0.01 % | 0.00 % |
| <b>Maximum</b> | 94.97 %  | 51.33 %     | 48.78 %         | 0.27 %                   | 0.01 % | 0.00 % |
| <b>Minimum</b> | 0.00 %   | 51.21 %     | 0.05 %          | 0.01 %                   | 0.00 % | 0.00 % |
| <b>StDev</b>   | 24.01 %  | 0.05 %      | 15.31 %         | 0.07 %                   | 0.00 % | 0 %    |
| <b>Avg/Max</b> | 0.48     | 1.00        | 0.19            | 0.29                     | 0.75   | 1      |

Figure 17: Histogram for task dependencies version

In figure 17 we can see that in some threads a lot of time is lost in some threads in synchronization state. We think this could be due to the tasks queue access.

We make the problem size 10 times bigger and check how it behaves. It is clear that we have the same issue here as represented in figure 18.

|                | Running  | Not created | Synchronization | Scheduling and Fork/Join | I/O    | Others |
|----------------|----------|-------------|-----------------|--------------------------|--------|--------|
| THREAD 1.1.1   | 97.96 %  | -           | 1.98 %          | 0.05 %                   | 0.01 % | 0.00 % |
| THREAD 1.1.2   | 53.46 %  | 44.19 %     | 2.30 %          | 0.03 %                   | 0.01 % | -      |
| THREAD 1.1.3   | 54.40 %  | 44.19 %     | 1.35 %          | 0.04 %                   | 0.01 % | -      |
| THREAD 1.1.4   | 54.53 %  | 44.19 %     | 1.23 %          | 0.04 %                   | 0.01 % | -      |
| THREAD 1.1.5   | 52.96 %  | 44.20 %     | 2.80 %          | 0.03 %                   | 0.01 % | -      |
| THREAD 1.1.6   | 43.23 %  | 44.19 %     | 12.54 %         | 0.03 %                   | 0.01 % | -      |
| THREAD 1.1.7   | 4.09 %   | 44.19 %     | 51.71 %         | 0.00 %                   | 0.01 % | -      |
| THREAD 1.1.8   | 55.24 %  | 44.19 %     | 0.52 %          | 0.04 %                   | 0.01 % | -      |
|                |          |             |                 |                          |        |        |
| <b>Total</b>   | 415.87 % | 309.36 %    | 74.44 %         | 0.25 %                   | 0.08 % | 0.00 % |
| <b>Average</b> | 51.98 %  | 44.19 %     | 9.30 %          | 0.03 %                   | 0.01 % | 0.00 % |
| <b>Maximum</b> | 97.96 %  | 44.20 %     | 51.71 %         | 0.05 %                   | 0.01 % | 0.00 % |
| <b>Minimum</b> | 4.09 %   | 44.19 %     | 0.52 %          | 0.00 %                   | 0.01 % | 0.00 % |
| <b>StDev</b>   | 23.75 %  | 0.00 %      | 16.43 %         | 0.01 %                   | 0.00 % | 0 %    |
| <b>Avg/Max</b> | 0.53     | 1.00        | 0.18            | 0.70                     | 0.78   | 1      |

Figure 18: Histogram for task dependencies version (10 times bigger problem size)

Now, we try to make the problem 100 times the size of the original, however when submitting, it seems like the program doesn't end and we can't obtain the corresponding trace and histogram. So we can't see whether or not it behaves the same way as with the original size or the 10 times the original.

## Optional:

Now we want to find the optimal values for `sort_size` and `merge_size`. First, we execute the script to find the optimal value for `sort_size`:

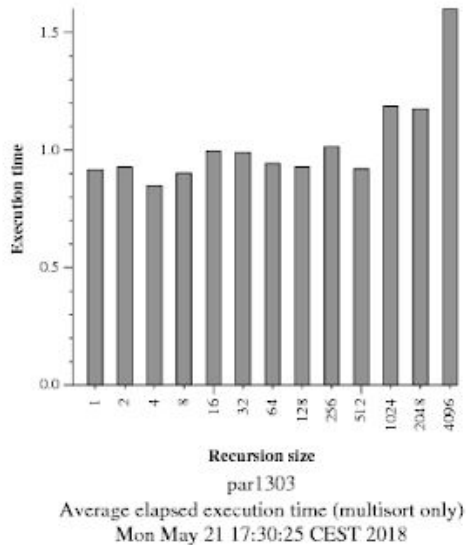



Figure 19: Plot of the execution times with different values of `sort_size`.

In the plot we can appreciate that seems there's no optimal number for `sort_size`. Almost all the execution times are the same, the difference is minimum. So maybe we can appoint that the `sort_size` doesn't affect in great measure the execution time, at least for this code.

To find the possible best or optimal `merge_size` we are going to try with `sort_size` = 4 and `sort_size` = 512. So we can appreciate the differences between a big number and a little one.

Now we want to find the optimal values for `merge_size`:

Figure 20:  
Plot of the execution times with `sort_size` = 4  
for different values of `merge_size`



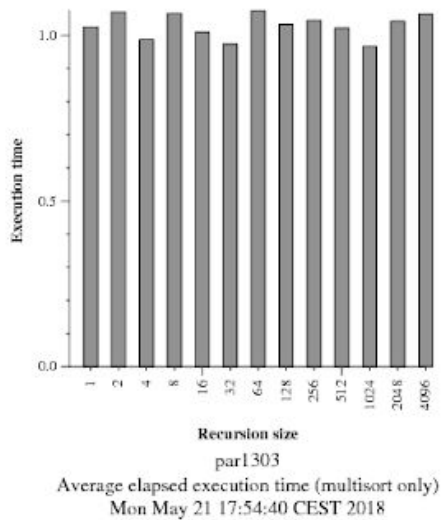


Figure 21:

Plot of the execution times with sort-size = 512 for different values of merge-size

With the results obtained we can see that there is not a remarkable difference that would allow us to assure that one of the options would be better than the rest of them. The execution time is very similar in all the executions. We would need a much more intense test of the program behaviour with the different values of merge-size and sort-size to see which are the most efficient combinations.

In this particular case, when comparing the execution times we can see that the best option seems to be sort-size = 4 and merge-size = 16, but, as said the difference is not remarkable enough.